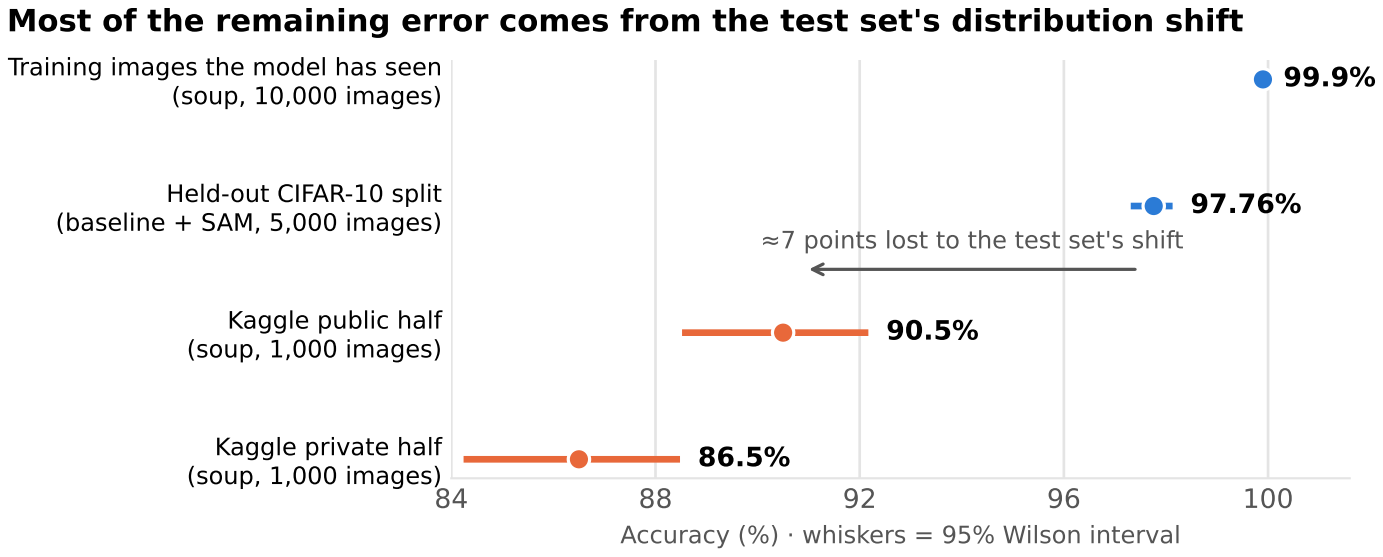
Most of the remaining error comes from the test set's distribution shift
Training images the model has seen
(soup, 10,000 images)
99.9%
Held-out CIFAR-10 split
(baseline + SAM, 5,000 images)
97.76%
≈7 points lost to the test set's shift
Kaggle public half
(soup, 1,000 images)
90.5%
Kaggle private half
(soup, 1,000 images)
86.5%
84
88
92
96
100
Accuracy (%) · whiskers = 95% Wilson interval